Computer Engineering Notes

Benjamin Davis [benjamin.j.davis96@gmail.com]

Contents

Chapter 1

Introduction

> **Definition 1.1: Example Box 2**
>
> This is an example 2 definition

Chapter 2

Secure Computer Systems

2.1   Introduction to Secure Computer Systems

There are threats to a cyber system. Specifically these can be summarized using the security mindset:

- **Threats** - Standard suite of bad-actors

- **Vulnerabilities** - Complexity of OS

- **Attacks** - Attacks against an OS are particularly attractive as they provide access to everything protected by the OS.

2.1.1   Trusted Computing Base: An Overview

An operating system is a set of software which makes it easier to use and share physical resources. These resources can range from Memory to CPU Cores to physical IO. The OS also manages these resources and can restrict access to them. In order to accomplish this an OS must have access to all the physical resources. Because of its complete access to all physical units the operating system provides access to all the physical components of the system and through this can also access all of the data processed by the system. In a secure system the OS can be considered/made to be a trusted computing base.

> **Definition 2.1: Trusted Computing Base**
>
> A software suite (usually and operating system) upon which other programs run that must be trusted for the security of said programs. (i.e. a software *BASE* which can be used and *TRUSTED* by other programs)

Trusted Computing Bases (TCBs) have several key functions, but the foremost is arguably that of a reference monitor. In it's role as a reference monitor a TCB maintains limited access to certain resources by the various programs which run on the TCB. That is to say that the TCB must monitor and control program references to these resources. It is critical that no reference can bypass a TCB, else the TCB loses its reference monitor capability and is defeated.

A properly implemented TCB has a handful of requirements:

1. **Tamper-Proof**

- Untrusted code cannot be able to operate the TCB

2. **Complete Mediation**

   - Untrusted applications must be unable to access protected resources without going through the TCB

   - The TCB cannot be bypassed by an untrusted program

3. **Correctness**

   - A TCB shall have no known vulnerabilities

   - A TCB vulnerability negates the previous two requirements.

Now it should be noted that a TCB does rely on hardware. Trust in hardware is its own challenge and beyond the scope of this chapter (see chapter 3 for more details). Thus for the remainder of this chapter it is prudent to assume the hardware is trusted.

[Ahamad, 2022]

2.1.2   TCB as a Reference Monitor

When acting as a reference monitor what tasks should the TCB perform?

1. **Authentication**

   - Who is the source of the request?

2. **Authorization**

   - Does the source of the request have permission to access the resource?

3. **Audit**

   - Can a admin check the above two requirements have been adhered to so far?

In order for a TCB to act as a reference monitor the TCB must be able to perform the above tasks in order to ensure trust.

[Ahamad, 2022]

2.1.3   Requirements for Trustworthiness

How do we decide to trust a TCB? Are there certain attributes which contribute to our trust of a TCB?

Generally trust in a TCB comes from two places. The first is what the TCB does. These are the claims about what the TCB is able to do and its reported limitations. The second is how well it does these things. It does not particularly matter how secure or trustworthy a TCB claims to be if it doesn't actually follow through on those promises. What the TCB does is pretty easy to judge, but how well is more difficult.

There are some things which can be analyzed to help support the second consideration. One of these is how the TCB is structured. Structuring allows for defense in depth and other strategies to be analyzed. Testing and formal verification are other strong metrics which allow

[Ahamad, 2022]

These cause other questions to arise. Can we truly trust any code which we have not written? It would seem no unless you are starting from the very base and writing everything yourself (of course even here it is impossible to fully trust the code as a single human is bound. to have made a mistake or two) Now one is tempted to reject this on the notion that if the source code can be reviewed for trojans or other "bugs" then the code can indeed be trusted. However, in practice this is not true as Ken Thompson pointed out in his life time achievement award. It is entirely possible to be using a tool which is built so as to re-insert a trojan when it rebuilds itself or introduce a different trojan when building any other programs. This problem forces complete trust to be derived only from a machine-code base up system which is entirely developed by yourself for complete trust. Alternatively, anyone who has touched the TCB or any part of the tool-chain to arrive at the TCB must be trusted implicitly. [Thompson, 1983]

2.1.4   TCSEC: The Orange Book

So if perfect trust is so hard to attain are there questions one can ask to determine whether a TCB can truly be trusted?

In the mid-1980s the DoD Security center published the *Trusted Computer System Evaluation Criteria*, colloquially known as "The Orange Book". This document puts forward a list of questions and requirements for certain levels of trust used by US Government secure computer systems. [DoD, 1985]

The Orange Book puts forward set of classes which evaluate the trust of a system as discussed below.

- **Class D**

    - Fails to meet the minimum requirements of TCB

- **Class C1**

    - Isolation of TCB

    - User Authentication

    - Access Control (discretionary)

- **Class C2**

    - All the features of C1

    - Accountability and Audit Requirements

    - Logging capability

- **Class B1**

    - All the features of C2

    - Well-defined TCB

    - Access Control (mandatory)

    - Penetration Testing

- **Class B2**

    - All the features of B1

    - Confinement and Covert Channels

– Well Structured TCB (such as modularity)

- **Class B3**

  – All the features of B2

  – Well defined security model

  – Separation of security code

  – Least Privilege

- **Class A1**

  – All the features of B3

  – Formal design verification

- **Class A2**

  – All the features of A1

  – Formal implementation verification

2.1.5   Introduction to Trusted Platform Module

Having software which exists at the TCSEC *A2* level is all well and good, but at some-point that TCB has to be loaded and booted. How do we know some bad actor has not changed the TCB and we are not booting a malicious or modified TCB? One option is the Trusted Platform Module (TPM).

> **Definition 2.2: Trusted Platform Module**
>
> A hardware "root" of trust which is included in a computer system.
>
> It performs an attestation of the operating system and platform by checking a digest of the TCB which is cryptographically strong.

In theory this solves the immediate issue since we are considering hardware trusted; however, TPMs are not without their own draw backs. TPMs allow hardware and OEM vendors to control what sorts of software (and OS) can be used on a specific hardware platform since other software would not pass the attestation of the TPM. This can lead to companies such as Apple locking users out of using their device except with authentic iOS or Apple signed software.

2.2   Design Principles for Secure Systems

Secure systems are complex and must accomplish a wide variety of tasks to be successful. If our task were to create an extremely simple system secure and trustworthy it would be relatively straight forward; however operating systems, and by extension TCB, are extremely complex. For a perfectly trustworthy system there should be no vulnerabilities or errors, since a single vulnerability is all that an adversary may need to break trust. So, how do we minimize the likelihood of security security weaknesses? The short answer is by following a key set of design principles.

2.2.1 Real-World Security

Physical security has been around far longer than secure computer systems, so let's start by looking at it for some basic principles. In the real-world only valuables are secured, and the more valuable the more secure we are likely to make the item. These valuables tend to be stolen for profit, though some items also may be targets for vandalism. The type of threat against the item determines the type of security which is necessary. Additionally, as we add more security to an item the types of threat actor which we encounter changes. For example banks are very secure, but also provide the opportunity for high-profit ventures, thus we can expect a smaller but more skilled group of thieves to continue to target the bank. Knowing this make up of threat actors and already implemented defenses allows a physical security firm to make better decisions about what new security features to add.

2.2.2 Design Principle 1: Security Cost

The first design principle which should be adhered to is the economics of security. Security cost must be commensurate with the treat level and asset value. Fundamentally, it makes sense that we should not spend more money trying to secure an item than that item is worth. Likewise, it is not necessary to make an item impossible to steal, instead it just needs to be expensive enough and risky enough that it is not worth it for a thief to try and steal it. But to go beyond that we also recognize that the more likely an item is to be targeted the more we need to spend to ensure its protection. From a defense standpoint this can be expressed mathematically. We would like to add defenses until:

$$C_D < V_D \text{ and } C_A > V_A \tag{2.1}$$

where C_D is the cost to the defender, V_D is the items value to the defender, C_A is the cost to the attacker, and V_A is the value to the attacker.

> **Design Principle 1: Security Cost**
>
> Security cost must be commensurate with the treat level and asset value.

2.2.3 Design Principle 2: User Acceptability

In addition to the economic cost of security there is another lest quantifiable cost: usability. Ideally security does not interfere with the usability of the system, but in practice this cannot be the case. For example password requirements and multi-factor authentication can greatly improve security, but they also add burden to a standard user.

In a perfect world all of these additional security measures would be handled automatically. However, systems are only secure as the weakest link. More and more often these days that weak link is the user. In order to remedy this issue additional requirements must be placed on the user, or at least offered to the user. The trade-off is that if there requirements are too cumbersome the user may opt not to use them or not to use the system that requires them. This trade-off is the user-acceptability design principle.

> **Design Principle 2: User Acceptability**
>
> The usability of a system must not be so compromised in the name of security that the intended users accept the requirements imposed on them.

2.2.4 Design Principle 3: System Complexity

Is it easier to analyze the security of a small program or a large one? The obvious answer is that smaller systems may be more easily analyzed than larger ones. But larger systems do provide the notable benefit of being more feature rich. Thus in a TCB it is necessary to balance these two trades. When possible fewer and simpler mechanisms should be used. However, this not always possible, so as systems get more complex more attention to the security of the systems needs to be heeded.

> **Design Principle 3: System Complexity**
>
> A secure system should be as simple as possible while still achieving the necessary functionality.

2.2.5 Design Principle 4: Least Privilege

When a user is running an application how do we decide what privileges that user and program should have to interact with resources? Should they have the bare minimum required to complete their task? Probably.

This is the principle of least privilege and it is the strongest principle from a security standpoint. However, in complex systems it can be tempting to bypass this system for a more convenient method of assigning privilege such as based on the user running the program, without regard for what the program actually needs. In fact this user-level privilege model is what Windows uses. Android on the other hand uses a stronger model much closer to least privilege as it grants permissions (i.e. privileges) based on what the application does in conjunction with what the user is willing to let it do. However, even this is not perfect least privilege since the program could request privileges it doesn't actually need if the app is a bad actor.

Ultimately, least privilege results in two key principles: separation of privileges and fail-safe default. Separation of privileges implies a finer grained access control. This might entail splitting files into different privileges. It may also include dividing different I/O devices in to different privilege requirements. Fail-safe defaults deny access when access is not explicitly granted. This means that some programs which are not granted a permission, but need it, could fail to execute properly. However, this is preferable to allowing malicious programs access to too much.

> **Design Principle 4: Least Privilege**
>
> A user and a program pair should be granted the least amount of privilege (i.e. access) to a system required in order to perform their approved task.

2.2.6   Design Principle 5: Defense in Depth

Can any defensive layer be perfect? No (and anyone who tells you differently is trying to sell you something). Instead we must assume each layer has a set of weaknesses which an attacker may exploit. Each of these layers may share some weaknesses, but they may also cover others. Thus if we combine two layers we can cover some of the weaknesses of the system and force a malicious force to use various types of attack which significantly increases the attack complexity. The more layers which are stacked on each other the more this complexity increases and the security increases. Through this multi-layer construction of security we gain additional defense through the depth (number of layers) of security.

Design Principle 5: Defense In Depth

A secure program should use multiple methods of defense in order to cover the vulnerabilities of one layer with the strength of another layer whenever possible.

Chapter 3

Hardware Trust

Chapter 4

Side-Channels

4.1 What is a Side-Channel?

4.2 Fine-Grained Analysis of Physically Observable SC

4.2.1 Frequency-Domain Analysis of EM Side-Channels

Remember that periodic activities have a spectral component at $f_1 = 1/t_1$ where t_1 is the time to complete the periodic activity. Tracking the code over time it will present with the spectral component of the current task. When another task is performed the spectral components will change which can be seen in the side channel and allows tracking of program flow.

Remember that a clock is a periodic activity on which the program execution can be modulated. This allows us to observe the spectrum at higher frequencies where there is less environmental noise.

Also remember that malware can be detected using the spectral components of the executed program. For example if malicious code is inserted during the loop the spectral component decrease in frequency compared to the expected EM frequency. However, if code is deleted during a loop the spectral component will increase and also be detectable. Finally, if code is inserted between loops it will be seen in the break between where the loops are expected to transition.

So this raises the question: how do we identify repetitive structures in the spectrum?

HDBSCAN

One powerful tool to automate the location of repetitive structures in the spectrum is known as HDBSCAN.

> **Definition 4.1: HDBSCAN**
>
> Hierarchical density-based spatial clustering of applications with noise

HDBSCAN works by:

1. Transform the space according to the density

2. Builds the minimum spanning tree of the distance weighted graphic

3. Constructs a cluster hierarchy of connected components

4. Condenses the cluster hierarchy based on the minimum cluster size

5. Extracted the stable clusters from the condensed tree

Once the clusters are identified peaks may be detected. The peaks correspond to the harmonics of the spectral component. These peaks are then used to check when the spectrum is in this loop and determine if the loop has been altered or left. If any of the harmonic peaks is different it is likely the loop has been altered or it is a different loop.

4.2.2   Dimension Reduction

An alternative to HDBSCAN is two-phase-dimension-reduction. As the name indicates two-phase-dimension-reduction works in two stages.

The first phase averages the magnitudes of the short-time fourier (STFT) transform using Equation 4.1. This enables a the noise and the number of samples to be reduced. Note that in this equation $X_l[k]$ is the STFT spectrogram's k^{th} component and N_R is the number of samples which are averaged together.

$$x_i[k] = \frac{1}{N_R} \sum_{l=1}^{N_R} |X_l[k]| \tag{4.1}$$

Phase two applies principal component analysis (PCA) to the result from phase one represented as a matrix (see Equation 4.2). In the matrix m represents the total number of measurements. $\hat{x}_i[k] = 10 \log_{10}(|x_i[k]|^2)$ is used in order to filter out the linear effects of the STFT. Thus each column in the matrix represents an averaged frequency component of the signal in the dB-domain.

$$\begin{bmatrix} \cdots & \hat{x}_1 & \cdots \\ \cdots & \hat{x}_2 & \cdots \\ & \vdots & \\ \cdots & \hat{x_m} & \cdots \end{bmatrix} \tag{4.2}$$

The first part of the PCA is then to find the model for the system. This is done through singular value decomposition (SVD). SVD provides a method of creating a series of orthonormal (un-correlated linear relationships) which can be used to reduce a large variable space into a significantly smaller variable space (i.e. can reduce the number of variables needed to represent a vector).

This is then used to create a smaller dimension array which can be easily analyzed using the k-th nearest neighbor algorithm (k-NN).

Using PCA provides several benefits. The first and foremost is that different operations use various frequency components, but tracking all frequency components can require a high-level of overhead. However, frequency components and corresponding thresholds must be conservatively set in order to ensure data size reduction without loss of variation between classes.

> **Algorithm 4.1: Two-Phase Dimension Reduction**
>
> 1. Measure Emanated Signal for T seconds
>
> 2. Reduce dimension using Phase One technique to get X
>
> 3. Apply Phase 2 to X to generate mode
>
> 4. Collect testing signal z and apply Phase One to get $\hat{z}$
>
> 5. Project the test signal $\mathcal{Z} = \hat{z} V_K \sigma_K^{-1}$
>
> 6. Apply k-NN algorithm to approximate the status of the device

4.2.3   Detecting Malware Using Frequency-Domain

One potential use for frequency-domain side-channels is to detect malware. This can be done by looking for anomalies in the spectrogram generated during program run. In order to easily find these it is important to understand the various spectrographic shapes which a loop can take.

Generally speaking there are three types of loop: fixed, control-flow, and nested. Each loop has a distinct spectral signature. Fixed loops are extremely stable as they perform the exact same operations repeatedly, thus their spectral signature is a single large peak. Control-flow loops on the other hand have internal branches which depend on the input to the loop. This gives these loops a multi-peak signature, with the different peaks representing the different paths through the loop. Finally nested loops contain many sub-loops. If these loops are adequately long they may be detected as multi-peak spectrums, but ordinarily they present as broad spectrum peaks.

In order to detect malware injection statistical tests are used. Due to the non-gaussian nature of the spectrograms it is necessary to use non-parametric tests. One such option is the K-S test which is sensitive to any difference between two groups of distribution.

Using a statistical test like this it is necessary to weight the tradeoff between detection latency and accuracy. On the one-hand we want to determine if malicious code is running on the processor as quickly as possible. However, we also do not want to flag good code or ignore malicious code. Thus we must increase the number of samples until the false rejection rate drops to a sufficient level.

How long this takes can depend heavily on the type of loop that is intended to be run. In general the fixed loop is the easiest to notice a change in due to its sharp singular feature, which is easily distinguished from other features. On the other end of the spectrum is the nested loop. Because of its broad peak many more samples are needed to distinguish generalized noise from a malicious attack.

4.2.4   Time-Domain Analysis of EM Side-Channels

While frequency-domain analysis is simple and extremely useful for determining broad information about the running program, it lacks the precision to dive deep into the inner workings of the program being observed. This is where Time-Domain analysis comes in. The time-domain provides a significantly finer grained analysis than frequency domain, which enables it to be used for program tracking from control flow down to individual instructions.

In order to achieve a clear signal for use in this sort of analysis the EM emanations must be extracted.

This is done through amplitude demodulation. Because the clock frequency is the strongest signal in the circuit its effects must be filtered out. To do so the EM signal is demodulated with respect to the clock frequency using Equation 4.3

$$x(t) = |r(t) \times e^{j2\pi f_c t}| \tag{4.3}$$

The first step towards using time-domain analysis is to determine what the expected emanations will look like. This is often done by running semi-random data through a program or part of program. If this is done enough times for each execution path an expected waveform which contains distinguishable characteristics can be determined from an average of traces. This allows a single trace, or set of traces averaged out, to be compared against the expected behavior for the different execution paths. By finding the correlation between the target signal and the expected signal the best and most probable match can be found.

This is a simple method, but can we do better?

Machine-Learning Techniques for EM Side-Channel Tracking

Machine learning techniques can be used to for automated fine-grain analysis of code execution. These techniques are not without issue though as some can lead to high computational explosion as the code-base increases in size. Machine learning is performed by first finding a set of training inputs which cover the entire flow-control space. These are then used on a known good training system to collect a series of waveforms which are used to train the ML model. A waveform can then be collected from a fielded system and analyzed in order to determine the flow path and the presence of any malicious code.

The program structure can be further leveraged to improve these techniques. By noting that some blocks of code will always execute together these can be used as singular references. It is also useful to note that these blocks will always process in a limited set of order (i.e. that not all code blocks can execute after any other code block). This allows an input wave form to be divided and analyzed using correlation to these subsets. A greedy tree-search can then be used to piece together the full execution path of the given wave form.

However, this tree search and correlation approach is limited in several key ways. Firstly it is difficult to scale to larger programs due to the infeasibility of determining all of the flow control paths and the additional memory and computational overhead of comparing a growing number of reference blocks. Additionally, the greedy tree-search is rather inefficient.

One way to improve this is to perform the path prediction with a hidden Markov model. This allows the algorithm to use additional states between within the blocks of code which have been identified. However, hidden Markov models still fail to scale well and rely on several key assumptions about the underlying program. Such assumptions includes statistically stationary transition probabilities and the presence of hidden information.

Speech Recognition Techniques for EM Side-Channel Analysis

An alternative to the relatively simple ML techniques previously discussed is the use of speech recognition techniques of the processing of a wave form. This option is believed to better manage the computational complexity of larger problems. The speech recognition used for EM analysis is based on forming a reduced dictionary and then using the dictionary to classify the signal being monitored.

In order to construct the original dictionary the EM signal is collected from a known-good source. This signal is then split into many overlapping short-duration fixed-length windows. Each windows is recorded as a dictionary entry (or a word in speech). A clustering technique can then be used to reduce the dictionary. One strong clustering candidate is threshold-based clustering which relies on continually growing centroids which aggregate surrounding dictionary entries into strong approximation models.

Once the reduced dictionary is constructed an EM signal is captured for analysis. Then each window in the EM signal is replaced by its best mach dictionary entry. If the dictionary entries were tagged with known code segments this allows for an approximate flow-control to be determined. Alternatively, if the desire is simply to monitor for anomalies then the error between the signal and the substituted entries can be used with a simple threshold to determine whether the variance falls into normal operating conditions. Increasing the complexity slightly to use a moving average filter for the anomaly threshold has been shown to substantially decrease the false positive rate.

Speech recognition techniques have excellent performance for malware detection without requiring a malware signature. Additionally, if malware detection is the only goal then the source code and control-flow graph are not required. However, as the program run-time increases the dictionary size may grow substantially. The speech recognition technique is also computationally inefficient during classification.

Neural Network Techniques for EM Side-Channel Analysis

Another method for fine-grained EM analysis is the use of a neural network. Many architectures of NN may be used for this technique. Multilayer perceptrons (MLP) have shown strong results in experiments. The neural network is tasked to find the expected amplitude of an EM signal at any given instance. In cases where real-time monitoring is not required a non-causal model may be used which significantly increases the efficacy of the method as it allows access to data before and after the instant being predicted.

While NNs require a longer training period than the other methods previously discussed they have some distinct advantages. A NN evaluation is constant time for a single evaluation which allows linear time evaluation of signals. Additionally, because the NN is a fixed size it has a constant memory complexity regardless of program size.

4.3   Covert-Channels

Covert-channels are a form of communication which is designed to not be obviously visible to the system. This often involves directly leveraging side-channels such as EM emanations in order to turn a processor into an RF transmitter. For example if a pair of loops can be found with different power characteristics they can be looped in a known way and create an amplitude modulated signal. Another style of covert channel could be to repeat specific instructions in order to enhance signal strength and increase the ease of detecting that instruction. Increasingly popular is the use of computer power management states. These have the benefit of being easily detectable, but tend to be slower than other alternatives.

4.3.1   Example Covert Channel

Here we demonstrate an example covert channel based on the idea of power states. Modern computers use demand based switching techniques to reduce power consumption when not under heavy load. These come in two types: performance states and processor states. Performance states refer to the current state of an active core, generally increasing the clock frequency when under heavy load and decreasing it when the

processor has enough overhead to handle the task with ease. Processor states on the other hand denote whether a processor is active or not. When a processor is parked (not active) it receives significantly lower currents and in some cases may be de-clocked entirely.

These states very clearly impact the emissions from the processor. However, they can impact more than the processor. When a CPU core is idle it may tell the voltage regulator module (VRM) to reduce the voltage level to the core. Because of the way modern voltage regulators operate this will significantly decrease the emissions from the device.

We can leverage these power states using a code which alternates between power states based on the bits we desire to send as shown in Algorithm 4.2.

Algorithm 4.2

```
1   void covert_channel(char* str) {
2     int dummy1;
3
4     // Step through chars
5     for (int i = 0; i < strLen(str); i++) {
6       // Step through char bits
7       for (int j = 0; j < 8; j++) {
8         if (bit_set(str[i], j)) {
9           // Keep Processor Active
10          for (int k=0; k < LOOP_PERIOD; k++)
11            dummy1 += dummy1 + 1;
12          // Put processor to sleep
13          usleep(SLEEP_PERIOD);
14        } else {
15          // Keep processor asleep
16          usleep(SLEEP_PERIOD * 2);
17        }
18      }
19    }
20  }
```

This algorithm works by shifting the power state from active to sleep for a 1, but leaving it asleep for 0 bits. Thus a RF receiver can see significant shifts in the magnitude of the frequency components.

An approach such as this is not without issues though. Consider the case where the program is interrupted by another activity such as the OS. This will cause the CPU to become active before returning control to the covert program. Additionally, in modern processors the power state of the CPU may not be strictly controlled by a single program as there are likely to be other programs which desire to run during the sleep time, or which run on the same physical processor using simultaneous multi-threading.

4.3.2   Modeling Covert Channels

Now that we've seen how to create a covert channel, we want to know how to model them.

First consider how the analog signals are created. The current in modern processors occurs from bit-flips

through transistors. However, modeling every transistor in a system is infeasible and unnecessary. Instead we find a corelation between the emissions and the processor's operating condition. Starting with a simple pipeline it has been shown that emanations vary depending on the state of the pipeline.

So it's simple then? Unfortunately not more needs to be done than just finding the power required for any given pipeline state. If we were to use a rectangular approximation of the signal with the power of any given state at the clock cycle we would get close, but still stray from the true output. As we see in subsection 4.3.2 a rectangular approximation is correct at the initial peak, but it quickly strays from the true signal.


Figure 4.1: Different Covert Channel Models

Intuitively the next option is model with a exponential curve. Again we run into a problem. While this is closer it misses the inherent oscillation which occurs as the signal ripples through the combinatorial circuits. Thus we must use a more complex model which includes some form of oscillation. The full final signal can be represented using:

$$y(t) = \sum_{n=1}^{\infty} x[n] \sin \left(\frac{2\pi(t - nT)}{T_0} \right) e^{-\theta(t-nT)} u(t - nT) \tag{4.4}$$

Where $x[n]$ is the initial power of the pipeline state and $u(t)$ is the unit step function.

$x[n]$ is found experimentally, but it must consider the entire state of the pipeline not just a single stage. However, this is known to be an aggregation of each individual stage so it can be found without an exponential search.

This model works great for pipeline, however we are missing some things. Namely we are missing micro-architectural events. Cache misses, branch predictions, and pipeline stalls are all going to play some substantial role on the covert channel. Thus any fully featured model ought to account for them.

4.4   Fault Injection Attacks

4.4.1   Types of Fault

Fault Effects

Fault Methods

4.4.2

4.5   Backscattering Side-Channels

Bibliography

[Ahamad, 2022] Ahamad, M. (2022). Lecture on secure computer systems, gatech cs 6238.

[DoD, 1985] DoD (1985). *Trusted Computer System Evaluation Criteria*. DoD Computer Security Center.

[Thompson, 1983] Thompson, K. (1983). Reflections on trusting trust. In *ACM Turing award lectures*. ACM.